AWS Certified Solutions Architect - Associate (SAA-C01)

Amazon AWS Certified Solutions Architect Associate Dumps Available Here at:

https://www.certification-questions.com/amazon-exam/aws-certified-solutions-architect-associate-dumps.html

Enrolling now you will get access to 272 questions in a unique set of AWS Certified Solutions Architect Associate dumps

Question 1

A Solutions Architect is designing the storage layer for a production relational database. The database will run on Amazon EC2. The database is accessed by an application that performs intensive reads and writes, so the database requires the LOWEST random I/O latency.
Which data storage method fulfills the above requirements?

Options:

A. Store data in a filesystem backed by Amazon Elastic File System (EFS).

B. Store data in Amazon S3 and use a third-party solution to expose Amazon S3 as a filesystem to the

database server.

C. Store data in Amazon Dynamo DB and emulate relational database semantics.

D. Stripe data across multiple Amazon EBS volumes using RAID 0.

Answer: D

Explanation:
When we perform the RAID 0 Striping of multiple volumes, IOPS are distributed among the volumes of a stripe. If you add another volume to RAID 0, you get the straight addition of IOPS throughput of that volume and additional volume size. Reference: https://cloudacademy.com/blog/amazon-aws-raid-0-configuration-on-ebs-volumes/

Question 2

A news organization plans to migrate their 20 TB video archive to AWS. The files are rarely accessed, but when they are, a request is made in advance and a 3 to 5-hour retrieval time frame is acceptable. However, when there is a breaking news story, the editors require access to archived footage within minutes.
Which storage solution meets the needs of this organization while providing the LOWEST cost of storage?

Options:

A. Store the archive in Amazon S3 Reduced Redundancy Storage.

B. Store the archive in Amazon Glacier and use standard retrieval for all content.

C. Store the archive in Amazon Glacier and pay the additional charge for expedited retrieval when needed.

D. Store the archive in Amazon S3 with a lifecycle policy to move this to S3 Infrequent Access after 30 days.

Answer: C

Explanation:
Expedited - Expedited retrievals allow you to quickly access your data when occasional urgent requests for a subset of archives are required. For all but the largest archives (250 MB+), data accessed using Expedited retrievals are typically made available within 1-5 minutes. Provisioned Capacity ensures that retrieval capacity for Expedited retrievals is available when you need it. For more information, see Provisioned Capacity.
Standard - Standard retrievals allow you to access any of your archives within several hours. Standard retrievals typically complete within 3-5 hours. This is the default option for retrieval requests that do not specify the retrieval option.
Bulk - Bulk retrievals are Glacier's lowest-cost retrieval option, which you can use to retrieve large amounts, even petabytes, of data inexpensively in a day. Bulk retrievals typically complete within 5-12 hours.

Question 3

The application tier for a workload runs on EC2 instances that are unevenly distributed across two Availability Zones. The instances are behind a Network Load Balancer and are accessed through layer 4 TCP connections.
The instances in the lesser populated Availability Zone are failing as the result of high CPU utilization.
Which configuration change can help mitigate the issue?

Options:

A. Modify the Network Load Balancer to enable sticky sessions

B. Modify the Network Load Balancer to enable cross-zone load balancing.

C. Switch to using an Application Load Balancer and enable sticky sessions.

D. Switch to using an Application Load Balancer and enable cross-zone load balancing.

Answer: D

Question 4

A company wants to migrate a three-tier web application to AWS. The company wants to control the placement of the instances and have visibility into underlying sockets and cores for licensing purposes.
Which compute model should a Solutions Architect choose to accomplish this task?

Options:

A. EC2 Reserved Instances

B. EC2 Spot Instances

C. EC2 Dedicated Hosts

D. EC2 Placement Groups

Answer: C

Question 5

A workload in an Amazon VPC consists of a single web server launched from a custom AMI. Session state is stored in a database.
How should the Solutions Architect modify this workload to be both highly available and scalable?

Options:

A. Create a launch configuration with a desired capacity of two web servers across multiple

Availability Zones. Create an Auto Scaling group with the AMI ID of the web server image. Use Amazon Route 53 latency-based routing to balance traffic across the Auto Scaling group.

B. Create a launch configuration with the AMI ID of the web server image. Create an Auto Scaling

group using the newly-created launch configuration, and a desired capacity of two web servers

across
multiple regions. Use an Application Load Balancer (ALB) to balance traffic across the Auto Scaling
group.

C. Create a launch configuration with the AMI ID of the web server image. Create an Auto Scaling

group using the newly-created launch configuration, and a desired capacity of two web servers across
multiple Availability Zones. Use an ALB to balance traffic across the Auto Scaling group.

D. Create a launch configuration with the AMI ID of the web server image. Create an Auto Scaling

group using the newly-created launch configuration, and a desired capacity of two web servers across
multiple Availability Zones. Use Amazon Route 53 weighted routing to balance traffic across the Auto
Scaling group.

Answer: C

Question 6

A Solutions Architect is trying to bring a data warehouse workload to an Amazon EC2 instance.
The data will reside in Amazon EBS volumes and full table scans will be executed frequently.
What type of Amazon EBS volume would be most suitable in this scenario?

Options:

A. Throughput Optimized HDD (st1)

B. Provisioned IOPS SSD (io1)

C. General Purpose SSD (gp2)

D. Cold HDD (sc1)

Answer: A

Explanation:
https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSVolumeTypes.html

Question 7

A company is developing several critical long-running applications hosted on Docker.
How should a Solutions Architect design a solution to meet the scalability and orchestration requirements on AWS?

Options:

A. Use Amazon ECS and Service Auto Scaling.

B. Use Spot Instances for orchestration and for scaling containers on existing Amazon EC2 instances.

C. Use AWS OpsWorks to launch containers in new Amazon EC2 instances.

D. Use Auto Scaling groups to launch containers on existing Amazon EC2 instances.

Answer: A

Explanation:
https://docs.aws.amazon.com/AmazonECS/latest/developerguide/service-auto-scaling.html

Question 8

A company with an existing AWS VPC is experiencing an increasing number of malicious attacks from a particular IP address range. The company wants to block all access from these IP addresses while the abuse patterns are being investigated.
How can access from the specified IPs be denied quickly and temporarily?

Options:

A. Use an AWS Marketplace solution to block access from the specified IP range.

B. Leverage NAT gateway on each instance to block access from the specified IP range.

C. Use network ACLs to block access from the specified IP range.

D. Create a rule in the security groups to block access from the specified IP range.

Answer: C

Question 9

A team must redesign the application layer of their SaaS Solution so that no matter request

volume required at any given time the solution will scale automatically to support the exact needs of that application. The team has decided to give serverless.
Which of the following service would best fit these requirements?

Options:

A. Amazon API Gateway and AWS Lambda

B. Amazon Cognito and Amazon DynamoDB

C. Amazon DynamoDB and AWS Code Deploy

D. AWS Lambda and an Application Load Balancer

Answer: A

Question 10

A Solutions Architect is designing a customer order processing application that will likely have high usage spikes.
What should the Architect do to ensure that customer orders are not lost before being written to an Amazon RDS database? (Choose two.)

Options:

A. Use Amazon CloudFront to deliver the application front end.

B. Use Elastic Load Balancing with a round-robin routing algorithm.

C. Have the orders written into an Amazon SQS queue.

D. Scale the number of processing nodes based on pending order volume.

E. Have a standby Amazon RDS instance in a separate Availability Zone.

Answer: C, D